

Question:

What **data structures** do you use the most and how do you choose between them? Illustrate with specific examples, avoiding vague generalities. Enthusiasm and strong conviction preferred!

Answer:

Throughout my academic and professional journey, I've consistently worked on a few key data structures that correspond with the complexity and nature of the problems I tackle.

- ❖ Arrays & N-dimensional Tensors: During my stint with the Image-Guided Medical Robotics Lab, handling high-dimensional data such as 3D optical coherence tomography images was paramount. For this, I habitually used arrays and tensors, leveraging their constant-time access property to quickly manipulate pixel values for real-time imaging and reconstruction. For instance, an adaptive imaging algorithm I crafted required assembling a 4D tensor to represent spatial-temporal relationships between different image slices.
- ❖ Hash Maps: While working on vision-based semantic segmentation with the Field Robotics Group, I relied heavily on hash maps (Python dictionaries), which are ideal for quick data retrieval without a sequential search. I mapped terrain features to their sensor data, which allowed constant-time complexity when fusing and accessing multimodal sensor data. This data structure was indispensable in ensuring that computational bottlenecks were minimized when analyzing vast streams of data in real-time.
- ❖ Graphs: At Mach Infinity Aerospace, my implementation of multi-layered graph neural networks for UAV state estimation illustrates the power of graphs. They represented the complex relationships and dynamic mapping required for 3D obstacle avoidance and navigation amongst sparsely depicted features, which could only be efficiently handled through a graph's connected nodes and edges.
- ❖ Dynamic Arrays (Vectors in C++): Confronted with uncertain amounts of data, such as when dealing with an unpredictable number of feature points in SLAM problems or varying list lengths in autonomous rover planning at the Robotics Institute, I often default to dynamic arrays. The ability of vectors to resize as needed underpins their utility in robotics, where environmental factors yield undetermined data quantities. Specifically, for Twilight SLAM, dynamic arrays were able to accommodate the unknown volume of image enhancers integrated into the pipeline without sacrificing speed or capacity.
- ❖ Sets: To ensure uniqueness in feature detection, especially during the Twilight SLAM project where repeatable feature identification was critical under varying illumination, I made use of sets extensively. Their properties automatically prevented duplicate feature descriptors, thus ensuring the integrity of the SLAM system's matching and tracking process.

Time complexity, space efficiency, and the ease of implementation are all factors I carefully balance from implementation viewpoint to ensure my solutions are not just workable but optimized for performance and scalability.

Question:

While C/C++ are famous for **pointers and the bugs associated with them**, understanding pointers is critical for effective programming in every language. Please provide a couple concrete examples of pointer-based bugs in a scripting language (e.g. Python or Javascript) and then also show the equivalent bug in C/C++ (or another system language).

Answer:

In my experience across various languages, understanding the intricacies of pointers and reference types is vital in avoiding common/critical bugs. In scripting languages like Python and JavaScript, which abstract away the explicit use of pointers, these issues often manifest in the form of reference errors.

Example 1:

Python: A scenario I encountered involved shared mutable default arguments in function definitions. For instance:

```
def append_to_list(element, shared_list=[]):
    shared_list.append(element)
    return shared_list
```

This seems innocent but leads to unexpected behavior:

```
list1 = append_to_list('a')
list2 = append_to_list('b')
# list1 and list2 both contain ['a', 'b'] which was not the intention
```

Here, `list1` and `list2` reference the same default list due to Python's handling of mutable default arguments, akin to unintended pointer aliasing.

C/C++: The essence of these bugs can be correlated to the misuse of pointers in C/C++:

```
void addToIntVector(int value, std::vector<int>* vec = nullptr) {
    static std::vector<int> defaultVec;
    if (vec == nullptr) {
        vec = &defaultVec;
    }
    vec->push_back(value);
}

// Later in the code
std::vector<int>* vec1 = nullptr;
std::vector<int>* vec2 = nullptr;
addToIntVector(1, vec1); // vec1 points to the static defaultVec
addToIntVector(2, vec2); // vec2 too points to now modified defaultVec
// vec1 and vec2 now point to a vector containing [1, 2] not two separate vectors
```

This happens because ``static`` in C/C++ creates persistence across function calls similar to the default argument list in Python. The pointers ``vec1`` and ``vec2`` both end up pointing to ``defaultVec``, leading to unexpected shared state modifications.

Example 2:

Python: A common pointer-based error in Python arises when I confuse deep and shallow copying. Shallow copying only copies the references to the objects, not the objects themselves:

```
import copy

original_list = [[1], [2], [3]]
shallow_copied_list = copy.copy(original_list)
shallow_copied_list[0][0] = 'a'

print(original_list)  # Output: [['a'], [2], [3]]
```

Here, modifying the shallow copy also affects the original because they contain references to the same nested objects. This type of bug is a form of aliasing issue due to the shared references.

C/C++: In C++, a similar mistake occurs when copying pointers. A shallow copy is generated when a pointer is copied, but not the content it points to:

```
#include <iostream>

class Container {
public:
    int* data;
    Container(int size) { data = new int[size]; }
    ~Container() { delete[] data; }
};

int main() {
    Container a(3);
    a.data[0] = 1; a.data[1] = 2; a.data[2] = 3;

    Container b = a;  // Shallow copy; b.data points to the same memory as a.data
    b.data[0] = 'a';

    std::cout << a.data[0] << std::endl;  // Output: 97 (ASCII value for 'a')

    return 0;
}
```

After the assignment ``Container b = a;``, both ``a.data`` and ``b.data`` point to the same memory region. Modifications through ``b`` will affect ``a``'s data as well, leading to aliasing.

Example 3:

Python: Another pointer-related issue in Python is unintended sharing of objects between instances:

```
class Accumulator:
    values = []

    def add(self, value):
        self.values.append(value)

acc1 = Accumulator()
acc2 = Accumulator()
acc1.add(10)

print(acc2.values)  # Output: [10]
```

This happens because `values` is a class attribute, shared by all instances of `Accumulator` due to its definition at the class level and not the instance level.

C/C++: In C++, a static class member has a parallel implication - shared ownership between instances:

```
#include <iostream>
#include <vector>

class Accumulator {
public:
    static std::vector<int> values;

    void add(int value) {
        values.push_back(value);
    }
};

std::vector<int> Accumulator::values;

int main() {
    Accumulator acc1, acc2;
    acc1.add(10);

    for (int val : acc2.values) {
        std::cout << val << " "; // Output: 10
    }
    return 0;
}
```

Both `acc1` and `acc2` share the same `values` vector because it is declared as `static`. Changes made by any instance will be seen by all other instances.

Question:

Gaia AI foresters perform data collection by hiking with our sensor backpack. Give an overview of a complete **factor graph localization** system for a sensor backpack in a forest environment. What specific low-level libraries and algorithms would you use? What kinds of factors would you use? What problems are likely to come up and how would you try to resolve them to make the localization robust?

Answer:

At the heart of the system, using a factor graph would allow us to integrate various sensor outputs into a coherent estimate of the hiker's trajectory and the backpack's position in the forest.

While designing the system I would use the following **low-level libraries**:

- ❖ GTSAM (Georgia Tech Smoothing and Mapping library) would be utilized to build and optimize the factor graph. It is well-suited for problems requiring non-linear optimization and has a rich set of tools for handling simultaneous localization and mapping (SLAM).
- ❖ PCL (Point Cloud Library) would be employed for processing and filtering dense LiDAR data. Algorithms within PCL can be used for feature extraction, segmentation, and object recognition—all critical in the context of SLAM and particularly challenging in a cluttered forest.
- ❖ OpenCV could be used for image processing tasks that complement LiDAR data, providing visual feature extraction that can be valuable in areas where GPS signals are unreliable.
- ❖ Eigen, high-level C++ library to use alongside GTSAM and PCL for efficient mathematical computations, which are imperative in optimization and transformation tasks.
- ❖ Boost, this collection of C++ libraries would be used to complement GTSAM and other libraries, filling in with data structures and algorithms not natively supported by them.

and **Algorithms**:

- ❖ iSAM2 (Incremental Smoothing and Mapping): Within the GTSAM library, iSAM2 stands out for its real-time SLAM capabilities, particularly useful for continuously updating factor graphs as new sensor data arrives.
- ❖ LOAM (Lidar Odometry and Mapping in Real-time): For processing LiDAR data, LOAM can provide high-precision, real-time odometry and mapping, which is advantageous when traversing complex forest terrain.
- ❖ DBoW2 (Bag of Words Library): For implementing loop closure detection in visual SLAM, this library helps identify previously visited locations for robustness against drift.

Factors I would use:

- ❖ Odometry Factors: These factors represent the motion model of the backpack based on data from inertial measurement units (IMUs) and/or step counters. (for relative motion estimates between successive poses)
- ❖ Measurement Factors (Sensor Integration): These factors encapsulate the direct measurements from the sensors, such as:
- ❖ LiDAR Factors: For the relative position estimation using point clouds and laser scan matching.
- ❖ Camera Factors: For feature-based constraints garnered from visual odometry or photogrammetry. They leverage data from the stereo or monocular cameras for visual SLAM.
- ❖ Environmental Factors: If there are reliable environmental maps or models (like digital elevation models (DEM)), factors derived from this data can be used to enhance the localization with a priori knowledge, like matching known tree positions or following known paths.

- ❖ **Loop Closure Factors:** Critical for long-term SLAM accuracy, these factors are based on recognizing a previously visited location using the backpack's sensors. This detection allows the system to minimize error accumulation over time by 'closing the loop' and adjusting the trajectory estimate.
- ❖ **Non-Motion Factors:** These represent non-dynamic constraints such as the assumed rigidity between sensors mounted on the backpack (for example, the fixed spatial relationship between a LiDAR and a camera), ensuring that different sensor data can be correctly correlative.
- ❖ **Absolute Position Factors (GPS):** Only If available in some open patches can be used as a soft prior for absolute position information to correct cumulative drift from odometry and sensor integration factors.

Likely problems and potential solutions:

- ❖ **Sensor Drift and Accumulative Errors:** Over time, small errors can compound leading to a significant drift. Leverage periodic 'reset' conditions, where high-confidence sensor readings (like very clear GPS signals) can help re-anchor the localization model. Incorporating IMUs with high bias stability and advancements like zero-velocity updates (ZUPT) can also reduce drift from inertial sensors.
- ❖ **Data Association Issues:** The system may incorrectly match sensor observations to the generated map due to ambiguities. Utilize sophisticated data association algorithms such as joint compatibility branch and bound (JCBB) or RANSAC to mitigate false positives in observation-to-map matching.
- ❖ **Dynamic vegetation and lighting:** Moving foliage due to wind, changes in lighting, and seasonal variations can affect sensor readings. Incorporate adaptive algorithms to differentiate between static and dynamic objects, possibly using machine learning for sparse scene reconstruction and/or adaptive illumination thresholding. Also, use robust feature extraction techniques in visual processing to minimize the impact of lighting changes.
- ❖ **Repetitive Patterns and Lack of Distinct Features:** Forests have recurring patterns that can confuse loop closure detection in SLAM. Implement more sophisticated loop closure techniques by combining local with global descriptors and using deep learning models to enhance feature distinctiveness for reliable recognition of previously visited locations.
- ❖ **Sensor Noise and Calibration Errors:** Inherent sensor inaccuracies and system calibration errors can degrade localization accuracy over time. Develop a continuous calibration routine where the system self-calibrates in real-time, reconciling sensor data with known physical models or periodic calibration waypoints. Also, include noise characterization of sensors in the factor graph, allowing the system to account for typical sensor errors.